

```
In [4]: import pandas as pd
        from nltk.corpus import stopwords
        import tensorflow as tf
        import numpy as np

        from tensorflow import keras
        from keras import models, layers

        from sklearn.preprocessing import LabelEncoder
        np.random.seed(1234)
```

```
In [12]: df = pd.read_csv('dataset.csv', encoding='latin1')
        df.pop('Language')

        print(df.head())
        df = df.truncate(before=0, after=20000)
        i = np.random.rand(len(df)) < 0.8
        train = df[i]
        test = df[~i]


        print("train data size: ", train.shape)
        print("test data size: ", test.shape)
```

	Text	Label
0	@Charlie_Corley @Kristine1G @amyklobuchar @Sty...	litigious
1	#BadBunny: Como dos gotas de agua: Joven se di...	negative
2	<a href="https://t.co/YJNi00p1JV">https://t.co/YJNi00p1JV</a> ( <a href="https://t.co/YJNi00p1JV">https://t.co/YJNi00p1JV</a> ) Flagstar Bank disclose...	litigious
3	Rwanda is set to host the headquarters of Unit...	positive
4	OOPS. I typed her name incorrectly (todayâ s ...	litigious

```
train data size: (16016, 2)
test data size: (3985, 2)
```

```
In [13]: import matplotlib.pyplot as plt
        df['Label'].value_counts().plot(kind='bar')
```

Out[13]: <AxesSubplot:>



```
In [14]: num_labels = 2
        vocab_size = 10000
        batch_size = 100

        tokenizer = keras.preprocessing.text.Tokenizer(num_words=vocab_size)
        tokenizer.fit_on_texts(train.Text)

        x_train = tokenizer.texts_to_matrix(train.Text, mode='tfidf')
        x_test = tokenizer.texts_to_matrix(test.Text, mode='tfidf')

        encoder = LabelEncoder()
        encoder.fit(train.Label)
        y_train = encoder.transform(train.Label)
        y_test = encoder.transform(test.Label)

        # check shape
        print("train shapes:", x_train.shape, y_train.shape)
        print("test shapes:", x_test.shape, y_test.shape)

        train shapes: (16016, 10000) (16016,)
        test shapes: (3985, 10000) (3985,)
```

```
In [15]: model = models.Sequential()
model.add(layers.Dense(16, activation='relu', input_shape=(10000,)))
model.add(layers.Dense(16, activation='relu'))
model.add(layers.Dense(1, activation='sigmoid'))
```

```
In [16]: model.compile(optimizer='rmsprop',
                      loss='binary_crossentropy',
                      metrics=['accuracy'],
                      run_eagerly=True)
```

```
In [17]: x_val = x_train[:10000]
partial_x_train = x_train[10000:]

y_val = y_train[:10000]
partial_y_train = y_train[10000:]
```

```
In [18]: history = model.fit(partial_x_train,
                             partial_y_train,
                             epochs=20,
                             batch_size=512,
                             validation_data=(x_val, y_val))
```

```
Epoch 1/20
12/12 [=====] - 1s 88ms/step - loss: -0.1674 - accuracy: 0.2902 - val_loss: -1.0230 - val_accuracy: 0.2790
Epoch 2/20
12/12 [=====] - 1s 60ms/step - loss: -1.7497 - accuracy: 0.2970 - val_loss: -2.4392 - val_accuracy: 0.2829
Epoch 3/20
12/12 [=====] - 1s 60ms/step - loss: -3.4697 - accuracy: 0.3072 - val_loss: -4.1710 - val_accuracy: 0.2879
Epoch 4/20
12/12 [=====] - 1s 54ms/step - loss: -5.6989 - accuracy: 0.3150 - val_loss: -6.3949 - val_accuracy: 0.2952
Epoch 5/20
12/12 [=====] - 1s 54ms/step - loss: -8.5119 - accuracy: 0.3216 - val_loss: -9.0999 - val_accuracy: 0.3018
Epoch 6/20
12/12 [=====] - 1s 55ms/step - loss: -11.8949 - accuracy: 0.3276 - val_loss: -12.2233 - val_accuracy: 0.3056
Epoch 7/20
12/12 [=====] - 1s 57ms/step - loss: -15.7667 - accuracy: 0.3311 - val_loss: -15.7360 - val_accuracy: 0.3073
Epoch 8/20
12/12 [=====] - 1s 62ms/step - loss: -20.1639 - accuracy: 0.3343 - val_loss: -19.7524 - val_accuracy: 0.3076
Epoch 9/20
12/12 [=====] - 1s 58ms/step - loss: -25.1353 - accuracy: 0.3356 - val_loss: -24.2254 - val_accuracy: 0.3084
Epoch 10/20
12/12 [=====] - 1s 62ms/step - loss: -30.6807 - accuracy: 0.3373 - val_loss: -29.1672 - val_accuracy: 0.3090
Epoch 11/20
12/12 [=====] - 1s 58ms/step - loss: -36.8734 - accuracy: 0.3384 - val_loss: -34.7358 - val_accuracy: 0.3096
Epoch 12/20
12/12 [=====] - 1s 58ms/step - loss: -43.7872 - accuracy: 0.3403 - val_loss: -40.7862 - val_accuracy: 0.3117
Epoch 13/20
12/12 [=====] - 1s 63ms/step - loss: -51.4105 - accuracy: 0.3421 - val_loss: -47.4346 - val_accuracy: 0.3124
Epoch 14/20
12/12 [=====] - 1s 59ms/step - loss: -59.7234 - accuracy: 0.3436 - val_loss: -54.7401 - val_accuracy: 0.3124
Epoch 15/20
12/12 [=====] - 1s 53ms/step - loss: -68.7914 - accuracy: 0.3438 - val_loss: -62.6921 - val_accuracy: 0.3129
Epoch 16/20
12/12 [=====] - 1s 55ms/step - loss: -78.7059 - accuracy: 0.3449 - val_loss: -71.2958 - val_accuracy: 0.3133
Epoch 17/20
12/12 [=====] - 1s 54ms/step - loss: -89.4085 - accuracy: 0.3447 - val_loss: -80.5155 - val_accuracy: 0.3135
Epoch 18/20
12/12 [=====] - 1s 55ms/step - loss: -100.8420 - accuracy: 0.3459 - val_loss: -90.3603 - val_accuracy: 0.3135
Epoch 19/20
12/12 [=====] - 1s 53ms/step - loss: -113.1405 - accuracy: 0.3471 - val_loss: -100.8734 - val_accuracy: 0.3134
Epoch 20/20
12/12 [=====] - 1s 52ms/step - loss: -126.2991 - accuracy: 0.3482 - val_loss: -112.2095 - val_accuracy: 0.3134
```

Took very little time, with 31% validation accuracy

```
In [19]: from sklearn.metrics import classification_report

pred = model.predict(x_test)
pred = [1.0 if p>= 0.5 else 0.0 for p in pred]
print(classification_report(y_test, pred))

125/125 [=====] - 0s 2ms/step
      precision    recall  f1-score   support

     0       0.97       0.15       0.27       833
     1       0.29       1.00       0.45      1123
     2       0.00       0.00       0.00      1108
     3       0.00       0.00       0.00       921

 accuracy          0.31      3985
 macro avg       0.32       0.29       0.18      3985
 weighted avg    0.28       0.31       0.18      3985
```

C:\Users\BillP\AppData\Local\Packages\PythonSoftwareFoundation.Python.3.9\_qbz5n2kfra8p0\LocalCache\local-packages\Python39\site-packages\sklearn\metrics\\_classification.py:1318: UndefinedMetricWarning: Precision and F-score are ill-defined and being set to 0.0 in labels with no predicted samples. Use `zero\_division` parameter to control this behavior.  
 \_warn\_prf(average, modifier, msg\_start, len(result))  
 C:\Users\BillP\AppData\Local\Packages\PythonSoftwareFoundation.Python.3.9\_qbz5n2kfra8p0\LocalCache\local-packages\Python39\site-packages\sklearn\metrics\\_classification.py:1318: UndefinedMetricWarning: Precision and F-score are ill-defined and being set to 0.0 in labels with no predicted samples. Use `zero\_division` parameter to control this behavior.  
 \_warn\_prf(average, modifier, msg\_start, len(result))  
 C:\Users\BillP\AppData\Local\Packages\PythonSoftwareFoundation.Python.3.9\_qbz5n2kfra8p0\LocalCache\local-packages\Python39\site-packages\sklearn\metrics\\_classification.py:1318: UndefinedMetricWarning: Precision and F-score are ill-defined and being set to 0.0 in labels with no predicted samples. Use `zero\_division` parameter to control this behavior.  
 \_warn\_prf(average, modifier, msg\_start, len(result))

```
In [20]: max_features = 10000
maxlen = 10000
batch_size = 32

model = models.Sequential()
model.add(layers.Embedding(max_features, 128, input_length=maxlen))
model.add(layers.Conv1D(32, 7, activation='relu'))
model.add(layers.MaxPooling1D(5))
model.add(layers.Conv1D(32, 7, activation='relu'))
model.add(layers.GlobalMaxPooling1D())
model.add(layers.Dense(1))
model.summary()
```

Model: "sequential\_2"

Layer (type)	Output Shape	Param #
=====		
embedding (Embedding)	(None, 10000, 128)	1280000
conv1d (Conv1D)	(None, 9994, 32)	28704
max_pooling1d (MaxPooling1D)	(None, 1998, 32)	0
conv1d_1 (Conv1D)	(None, 1992, 32)	7200
global_max_pooling1d (GlobalMaxPooling1D)	(None, 32)	0
dense_6 (Dense)	(None, 1)	33
=====		
Total params: 1,315,937		
Trainable params: 1,315,937		
Non-trainable params: 0		

```
In [21]: model.compile(optimizer=tf.keras.optimizers.RMSprop(lr=1e-4), # set Learning rate
                      loss='binary_crossentropy',
                      metrics=['accuracy'])
```

C:\Users\BillP\AppData\Local\Packages\PythonSoftwareFoundation.Python.3.9\_qbz5n2kfra8p0\LocalCache\local-packages\Python39\site-packages\keras\optimizers\optimizer\_v2\rmsprop.py:140: UserWarning: The `lr` argument is deprecated, use `learning\_rate` instead.  
 super().\_\_init\_\_(name, \*\*kwargs)

```
In [22]: history = model.fit(partial_x_train,
                             partial_y_train,
                             epochs=10,
                             batch_size=512,
                             validation_split=0.2)
```

```
Epoch 1/10
10/10 [=====] - 146s 14s/step - loss: 3.5593 - accuracy: 0.2136 - val_loss: 2.4577 - val_accuracy: 0.2434
Epoch 2/10
10/10 [=====] - 136s 14s/step - loss: 2.1833 - accuracy: 0.2136 - val_loss: 1.6489 - val_accuracy: 0.2434
Epoch 3/10
10/10 [=====] - 129s 13s/step - loss: 1.4341 - accuracy: 0.2136 - val_loss: 1.0469 - val_accuracy: 0.2434
Epoch 4/10
10/10 [=====] - 132s 13s/step - loss: 0.8161 - accuracy: 0.2363 - val_loss: 0.4833 - val_accuracy: 0.2924
Epoch 5/10
10/10 [=====] - 133s 13s/step - loss: 0.1618 - accuracy: 0.2872 - val_loss: -0.2465 - val_accuracy: 0.2924
Epoch 6/10
10/10 [=====] - 137s 14s/step - loss: -2.3149 - accuracy: 0.2874 - val_loss: -6.5353 - val_accuracy: 0.2924
Epoch 7/10
10/10 [=====] - 135s 13s/step - loss: -7.6483 - accuracy: 0.2874 - val_loss: -6.5354 - val_accuracy: 0.2924
Epoch 8/10
10/10 [=====] - 137s 14s/step - loss: -7.6531 - accuracy: 0.2874 - val_loss: -6.5354 - val_accuracy: 0.2924
Epoch 9/10
10/10 [=====] - 135s 13s/step - loss: -7.6531 - accuracy: 0.2874 - val_loss: -6.5354 - val_accuracy: 0.2924
Epoch 10/10
10/10 [=====] - 137s 14s/step - loss: -7.6531 - accuracy: 0.2874 - val_loss: -6.5354 - val_accuracy: 0.2924
```

Took about 22 minutes, with 29% validation accuracy

```
In [23]: from sklearn.metrics import classification_report
```

```
pred = model.predict(x_test)
pred = [1.0 if p>= 0.5 else 0.0 for p in pred]
print(classification_report(y_test, pred))
```

```
125/125 [=====] - 16s 129ms/step
```

	precision	recall	f1-score	support
0	0.00	0.00	0.00	833
1	0.28	1.00	0.44	1123
2	0.00	0.00	0.00	1108
3	0.00	0.00	0.00	921
accuracy			0.28	3985
macro avg	0.07	0.25	0.11	3985
weighted avg	0.08	0.28	0.12	3985

C:\Users\BillP\AppData\Local\Packages\PythonSoftwareFoundation.Python.3.9\_qbz5n2kfra8p0\LocalCache\local-packages\Python39\site-packages\sklearn\metrics\\_classification.py:1318: UndefinedMetricWarning: Precision and F-score are ill-defined and being set to 0.0 in labels with no predicted samples. Use `zero\_division` parameter to control this behavior.  
\_warn\_prf(average, modifier, msg\_start, len(result))

C:\Users\BillP\AppData\Local\Packages\PythonSoftwareFoundation.Python.3.9\_qbz5n2kfra8p0\LocalCache\local-packages\Python39\site-packages\sklearn\metrics\\_classification.py:1318: UndefinedMetricWarning: Precision and F-score are ill-defined and being set to 0.0 in labels with no predicted samples. Use `zero\_division` parameter to control this behavior.  
\_warn\_prf(average, modifier, msg\_start, len(result))

C:\Users\BillP\AppData\Local\Packages\PythonSoftwareFoundation.Python.3.9\_qbz5n2kfra8p0\LocalCache\local-packages\Python39\site-packages\sklearn\metrics\\_classification.py:1318: UndefinedMetricWarning: Precision and F-score are ill-defined and being set to 0.0 in labels with no predicted samples. Use `zero\_division` parameter to control this behavior.  
\_warn\_prf(average, modifier, msg\_start, len(result))

```
In [24]: model = models.Sequential()
model.add(layers.Embedding(max_features, 8, input_length=maxlen))
model.add(layers.Flatten())
model.add(layers.Dense(16, activation='relu'))
model.add(layers.Dense(1, activation='sigmoid'))

model.compile(optimizer='rmsprop', loss='binary_crossentropy', metrics=['acc'])
model.summary()

history = model.fit(partial x train, partial y train, epochs=10, batch size=512, validation split=0.2)
Model: "sequential_3"
```

Layer (type)	Output Shape	Param #
embedding_1 (Embedding)	(None, 10000, 8)	80000
flatten (Flatten)	(None, 80000)	0
dense_7 (Dense)	(None, 16)	1280016
dense_8 (Dense)	(None, 1)	17

```
=====
Total params: 1,360,033
Trainable params: 1,360,033
Non-trainable params: 0
```

```
Epoch 1/10
5/5 [=====] - 3s 497ms/step - loss: -14.7516 - acc: 0.2874 - val_loss: -34.3933 - val_acc: 0.2924
Epoch 2/10
5/5 [=====] - 2s 452ms/step - loss: -60.2966 - acc: 0.2874 - val_loss: -81.0130 - val_acc: 0.2924
Epoch 3/10
5/5 [=====] - 2s 472ms/step - loss: -119.2836 - acc: 0.2874 - val_loss: -137.4727 - val_acc: 0.2924
Epoch 4/10
5/5 [=====] - 2s 456ms/step - loss: -189.7151 - acc: 0.2874 - val_loss: -203.0400 - val_acc: 0.2924
Epoch 5/10
5/5 [=====] - 2s 460ms/step - loss: -270.1909 - acc: 0.2874 - val_loss: -276.8249 - val_acc: 0.2924
Epoch 6/10
5/5 [=====] - 2s 460ms/step - loss: -360.1279 - acc: 0.2874 - val_loss: -358.9992 - val_acc: 0.2924
Epoch 7/10
5/5 [=====] - 2s 451ms/step - loss: -461.1561 - acc: 0.2874 - val_loss: -451.3604 - val_acc: 0.2924
Epoch 8/10
5/5 [=====] - 2s 436ms/step - loss: -573.3301 - acc: 0.2874 - val_loss: -552.8761 - val_acc: 0.2924
Epoch 9/10
5/5 [=====] - 2s 463ms/step - loss: -696.4377 - acc: 0.2874 - val_loss: -663.4859 - val_acc: 0.2924
Epoch 10/10
5/5 [=====] - 2s 454ms/step - loss: -830.0669 - acc: 0.2874 - val_loss: -783.5403 - val_acc: 0.2924
```

```
In [*]: from sklearn.metrics import classification_report
```

```
pred = model.predict(x_test)
pred = [1.0 if p>= 0.5 else 0.0 for p in pred]
print(classification_report(y_test, pred))
```

took about 30 seconds with 28% validation accuracy

Overall decently low accuracy, but data probably could have been handled better to suit the models.